



DATA

IMDb

FILM MEETS FIGURES

Cua, Lander Peter E.
Gaba, Jacob Bryan B.
Gonzales, Mark Edward M.
Lee, Hylene Jules G.



01

DATABASE SCHEMA

Sources, schema, and
denormalization strategy

02

ETL PIPELINE

Pipeline processes
and utilized strategies

03

OPTIMIZATION AND RESULTS

MySQL scripts, findings,
and query execution plan

04

TESTING

Metadata, content, and
preprocessing validation

05

CONCLUSION AND INSIGHTS

Summary of findings
and analyses

06

DASHBOARD DEMO

Data visualization
via Tableau



Database Sources IMDb IJS

directors_genres	
director_id	int
genre	varchar
prob	float

movies_directors	
director_id	int
movie_id	int

movies_genres	
movie_id	int
genre	varchar

roles	
actor_id	int
movie_id	int
role	varchar

directors	
id	int
first_name	varchar
last_name	varchar

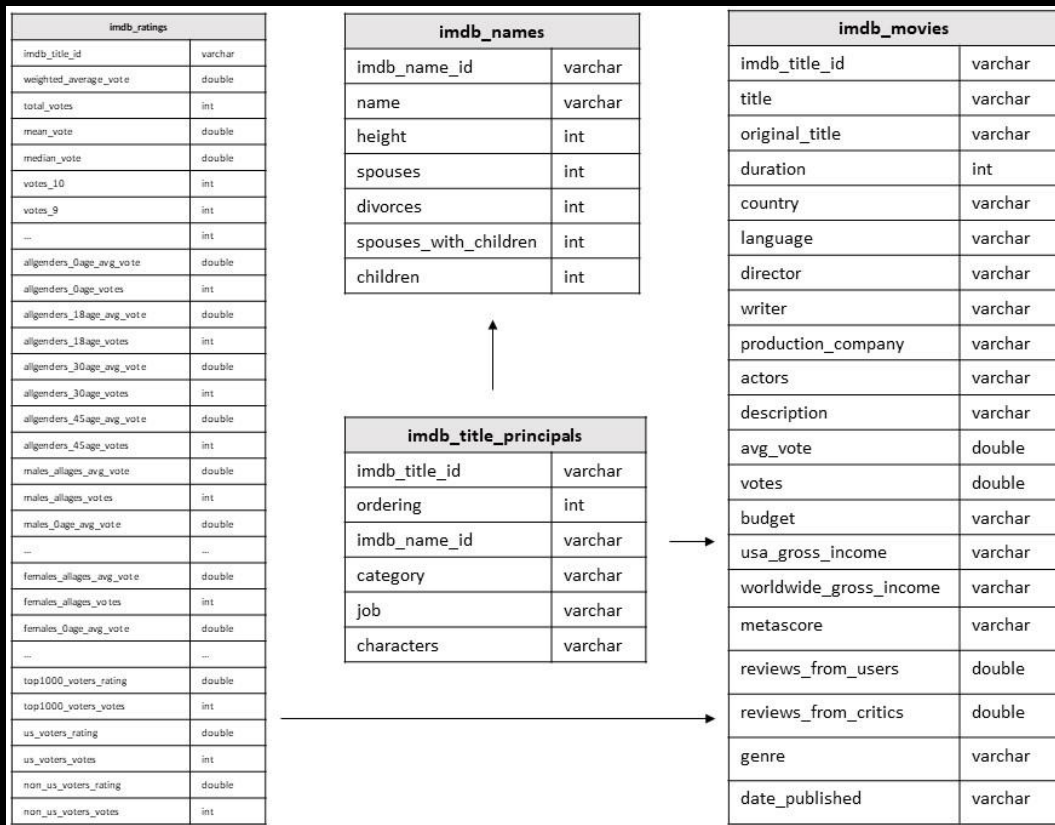
movies	
id	int
name	varchar
year	int
rank	float

actors	
id	int
first_name	varchar
last_name	varchar
gender	char

DATABASE SCHEMA



Database Sources IMDb Supplementary



DATABASE SCHEMA



DATA WAREHOUSE

- ✓ 1 fact table
- ✓ 7 dimension tables
- ✓ Star schema

ratings	
rating_id	int
total_num_votes	int
votes_10	int
votes_9	int
...	int
allgenders_0age_avg_vote	double
allgenders_0age_votes	int
allgenders_18age_avg_vote	double
allgenders_18age_votes	int
allgenders_30age_avg_vote	double
allgenders_30age_votes	int
allgenders_45age_avg_vote	double
allgenders_45age_votes	int
males_allages_avg_vote	double
males_allages_votes	int
males_0age_avg_vote	double
...	...
females_allages_avg_vote	double
females_allages_votes	int
females_0age_avg_vote	double
...	...
top1000_voters_rating	double
top1000_voters_votes	int
us_voters_rating	double
us_voters_votes	int
non_us_voters_rating	double
non_us_voters_votes	int

directors	
director_id	int
first_name	varchar
last_name	varchar
rate	double
gross	decimal

movies	
movie_id	int
name	varchar
year	int
rank	float
other_name	varchar
duration	int
language	varchar
production_company	varchar
description	varchar
votes	double
budget	varchar
usa_gross_income	varchar
worldwide_gross_income	varchar
metascore	varchar
reviews_from_users	double
reviews_from_critics	double

fact_table	
movie_id	int
country_id	int
genre_id	int
director_id	int
actor_id	int
role_id	int
rating_id	int

actors	
actor_id	int
first_name	varchar
last_name	varchar
gender	char

genre	
genre_id	int
genre	varchar

roles	
role_id	int
role	varchar

countries	
country_id	int
country	text

STAR SCHEMA JOINING TABLES

$$\Theta\left(\prod_{i=1}^m r_i\right)$$

Normalized

$$O((\max r_i)^m)$$

$$\Theta\left(\prod_{i=1}^p s_i\right)$$

Denormalized

$$O((\max s_i)^p)$$

$$m > p$$

Denormalization Strategies (Han, Kamber & Pei, 2012)

COLLAPSING



Applied to **movies**, **actors**, and **directors**

DIRECTOR DENORMALIZATION

Normalized:

```
movies(movie_id, name, ...)  
directors(director_id, first_name, last_name)  
movies_directors (director_id, movie_id)
```

Denormalized:

```
fact_table(movie_id, directors_id)  
directors(director_id, first_name, last_name)
```



Collapsing strategy with **auto-incrementing surrogate keys** applied to **ratings**

RATING DENORMALIZATION

Normalized:

```
movies(movie_id, name, ...)  
ijs2supplementary(movie_id, imdb_title_id)  
imdb_ratings(imdb_title_id, votes_10,  
votes_9, ...)
```

Denormalized:

```
fact_table(movie_id, country_id, genre_id,  
directors_id, actors_id, role_id, rating_id)  
ratings(ratings_id, votes_10, votes_9, ...)
```



VERTICAL PARTITIONING

Applied to **roles**, **genres**, and **countries**

- Makes up for the loss of conceptual distinction
- Optimization:

Separate long strings to another table to decrease number of pages

COUNTRY DENORMALIZATION

Normalized:

```
movies(movie_id, name, ..., country, ...)
```

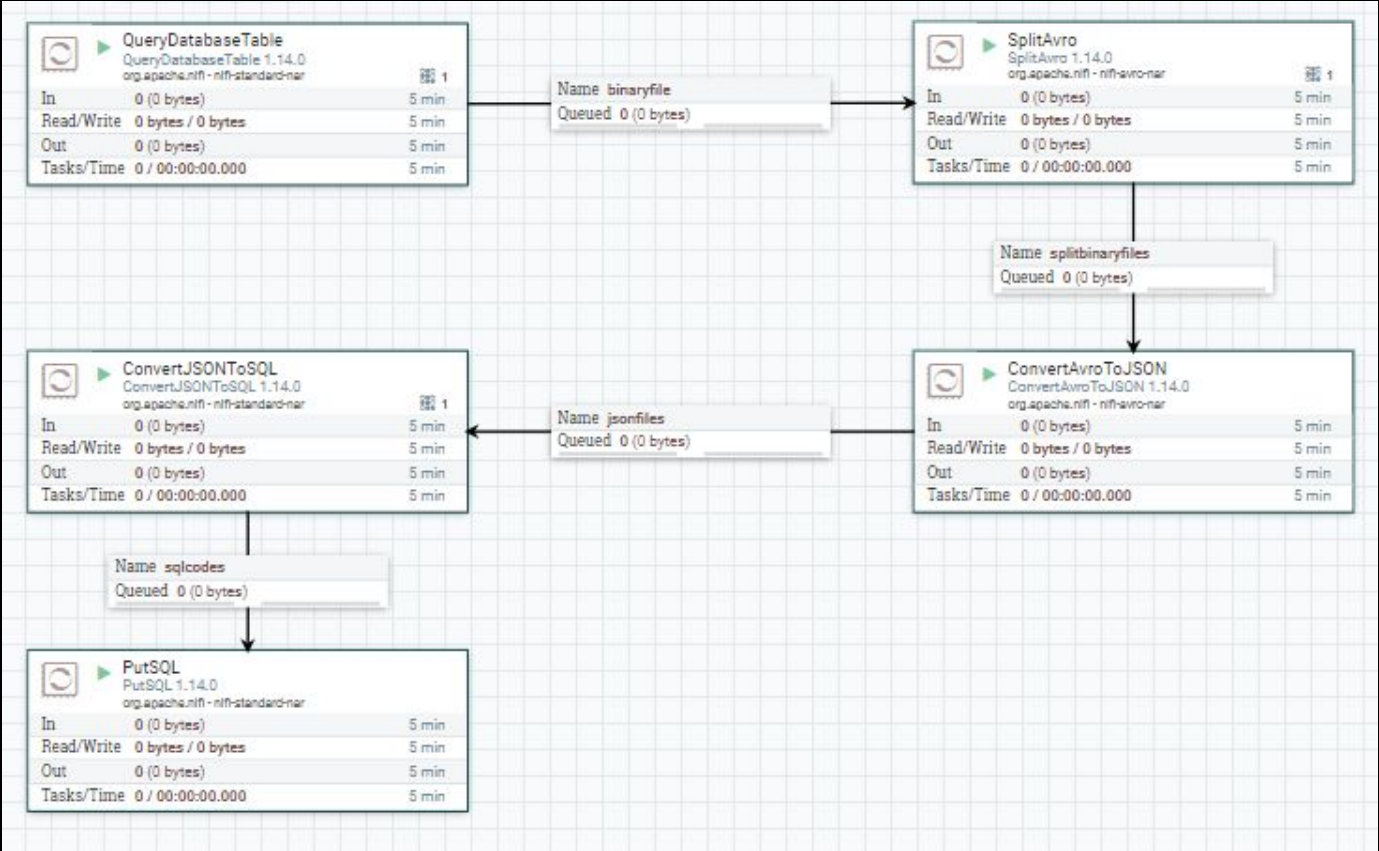
Denormalized:

```
fact_table(movie_id, country_id, genre_id,  
            directors_id, actors_id, role_id)
```

```
movies(movie_id, name, ...)
```

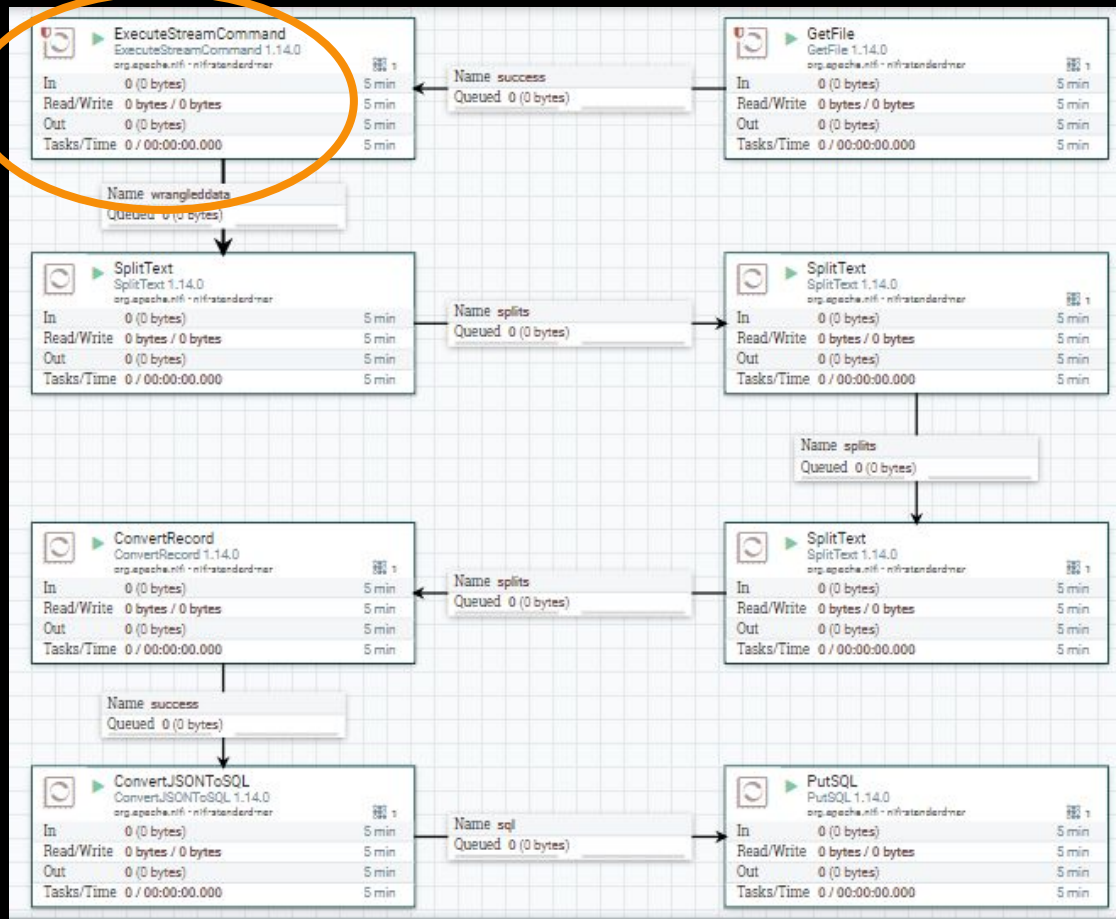
```
country(country_id, country)
```


Creating Local Copies IMDb IJS

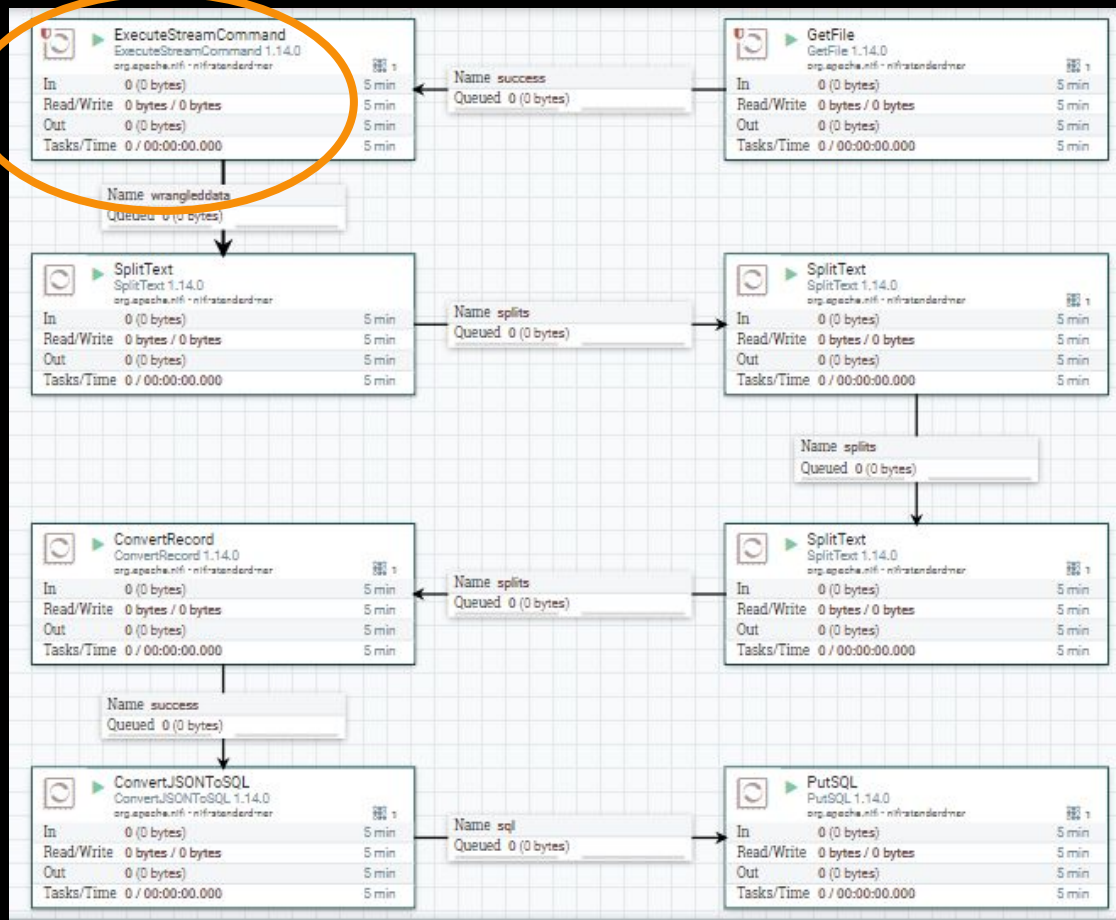


Creating Local Copies + Preliminary Data Wrangling

IMDb Supplementary



IMDb Supplementary



DATE STANDARDIZATION

Non-standardized:

10-09-1894

26/12/1906

1913

Standardized:

1894-09-10

1906-12-26

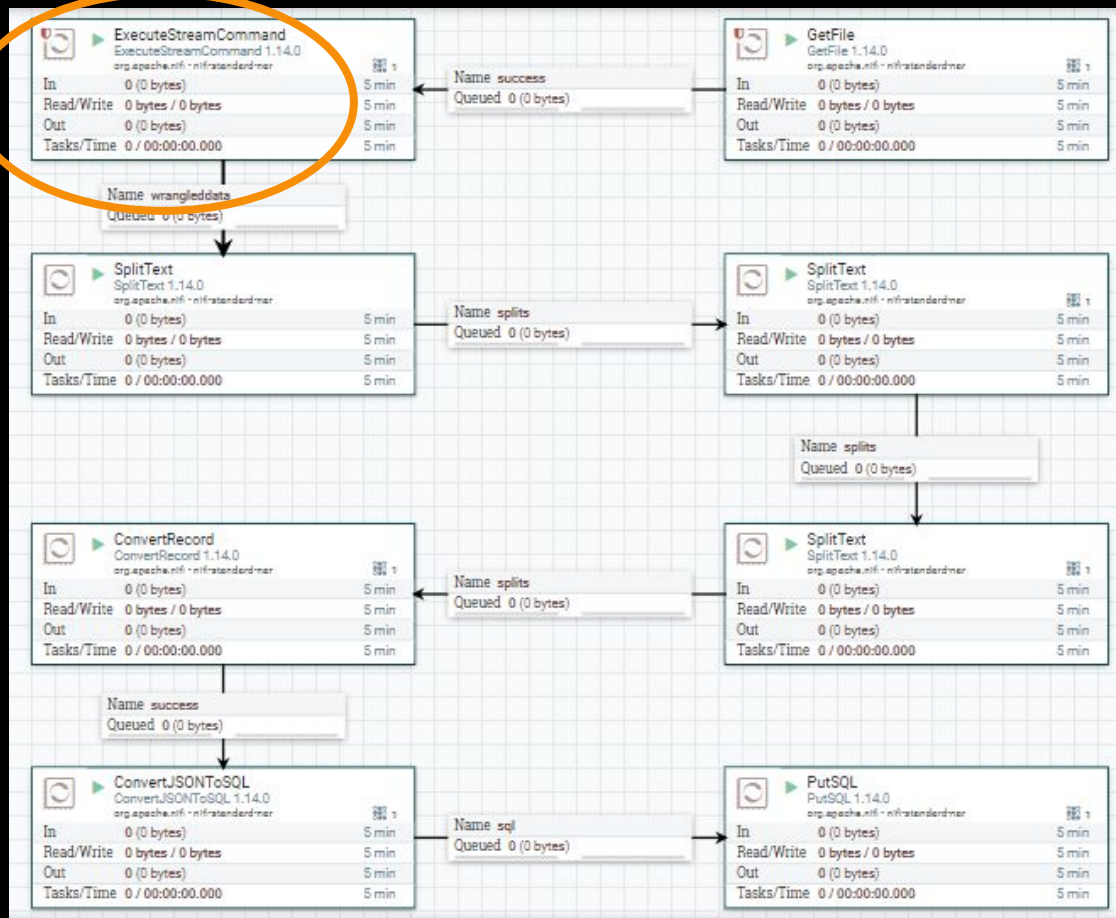
1913-01-01

ETL PIPELINE



Creating Local Copies + Preliminary Data Wrangling

IMDb Supplementary



DATE STANDARDIZATION

Non-standardized:

10-09-1894

26/12/1906

1913

Standardized:

1894-09-10

1906-12-26

1913-01-01

COUNTRY STANDARDIZATION

Non-standardized:

France, Germany, Japan

Exploded into separate rows:

France

Germany

Japan

ETL PIPELINE



Standardizing Movie Titles

Godfather, The (1972) → The Godfather



Standardizing Movie Titles

Godfather, The (1972) → The Godfather

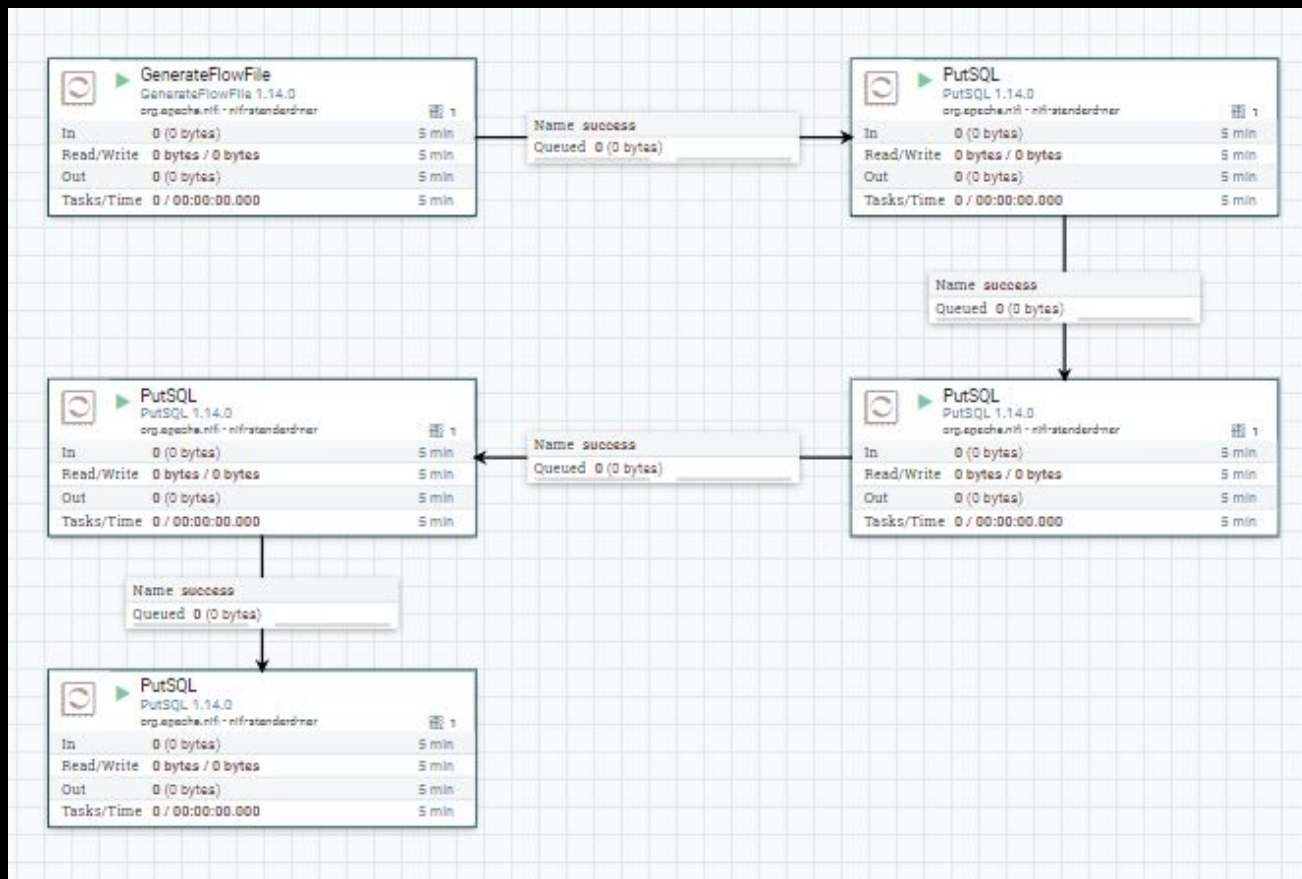
```
UPDATE movies SET `name` =  
RTRIM(REVERSE(SUBSTRING(REVERSE(`name`),  
LOCATE(" ",REVERSE(`name`)))))  
WHERE `name` IN (  
  SELECT `name`  
  FROM (  
    SELECT `name`  
    FROM movies  
    WHERE name LIKE "%(%)"  
  ) AS tmp  
);
```

```
UPDATE movies SET `name` =  
RTRIM(REVERSE(SUBSTRING(REVERSE(`name`),  
LOCATE(" ",REVERSE(`name`)))))  
WHERE `name` IN (  
  SELECT `name`  
  FROM (  
    SELECT `name`  
    FROM movies  
    WHERE name LIKE "%, The"  
  ) AS tmp  
);
```

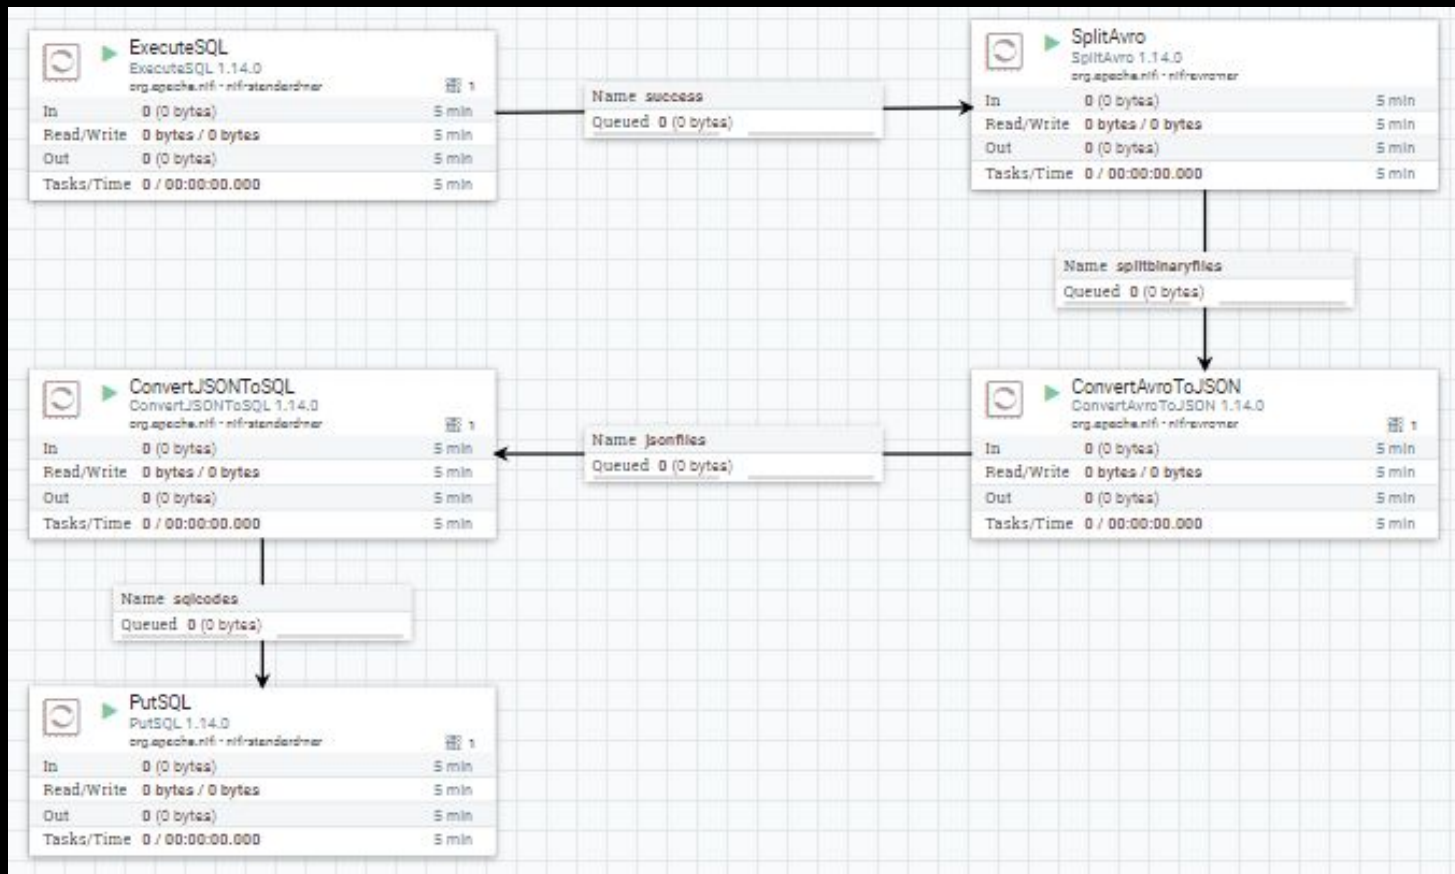
```
UPDATE movies SET `name` = CONCAT("The ",  
  `name`)  
WHERE `name` IN (  
  SELECT `name`  
  FROM (  
    SELECT `name`  
    FROM movies  
    WHERE name LIKE "%,"  
  ) AS tmp  
);
```

```
UPDATE movies SET `name` = SUBSTRING(`name`,  
  1, CHAR_LENGTH(`name`) - 1)  
WHERE `name` IN (  
  SELECT `name`  
  FROM (  
    SELECT `name`  
    FROM movies  
    WHERE name LIKE "%,"  
  ) AS tmp  
);
```





Transferring to the Data Warehouse



Constructing the Fact Table

```
SELECT movie_id, country_id, genre_id,  
       director_id, actor_id, role_id, rating_id  
FROM movies  
LEFT JOIN (movies_directors  
JOIN directors  
  ON directors.director_id =  
     movies_directors.director_id)  
  ON movies.movie_id =  
     movies_directors.movie_id  
LEFT JOIN (movies_actors  
JOIN actors  
  ON actors.actor_id =  
     movies_actors.actor_id)  
  ON movies.movie_id =  
     movies_actors.movie_id
```

```
LEFT JOIN genres  
  ON movies.movie_id =  
     movies_genres.movie_id  
LEFT JOIN roles  
  ON movies.movie_id = roles.movie_id  
  AND actors.actor_id = roles.actor_id  
LEFT JOIN (ijs2supplementary  
JOIN imdb_movies  
  ON ijs2supplementary.imdb_title_id =  
     imdb_movies.imdb_title_id)  
JOIN imdb_ratings  
  ON ijs2supplementary.imdb_title_id =  
     imdb_ratings.imdb_title_id)  
  ON movies.movie_id =  
     ijs2supplementary.movie_id;
```



Data Warehouse Size

Name	Rows
fact_table	6,617,850
actors	815,842
countries	184
directors	85,731
genres	21
movies	381,762
ratings	35,428
roles	2,513,767

▶ MACHINE SPECIFICATIONS

- ✓ **Processor:** AMD Ryzen 5 5600x 6-Core Processor
- ✓ **Processor Base Frequency:** 3.70 GHz
- ✓ **Memory:** 16GB DDR4 2667MHz
- ✓ **Disk:** 1TB Hard Disk Drive

▶ EXPERIMENTS

First Version normalized schema

Second Version without indexes

Third Version with indexes on attributes involved in join
operations

Fourth Version with composite indexes on attributes involved in
join, group by, and order by

Fifth Version revised, more optimized, equivalent queries (if possible)



QUERY PROCESSING

Q1. For each genre, what is the total number of movie votes per year and decade?

```
SELECT CONCAT(FLOOR(m.year/10) * 10, 's') AS
  decade_released, m.year, g.genre,
  (SUM(r.votes_10) + SUM(r.votes_9) +
  SUM(r.votes_8) + SUM(r.votes_7) +
  SUM(r.votes_6) + SUM(r.votes_5) +
  SUM(r.votes_4) + SUM(r.votes_3) +
  SUM(r.votes_2) + SUM(r.votes_1)) AS
  total_num_votes,
  SUM(r.votes_10) AS votes_10, SUM(r.votes_9)
  AS votes_9, SUM(r.votes_8) AS votes_8,
  SUM(r.votes_7) AS votes_7, SUM(r.votes_6) AS
  votes_6, SUM(r.votes_5) AS votes_5,
  SUM(r.votes_4) AS votes_4, SUM(r.votes_3) AS
  votes_3, SUM(r.votes_2) AS votes_2,
  SUM(r.votes_1) AS votes_1
FROM movies m
JOIN genres g ON g.movie_id = m.movie_id
JOIN ijs2supplementary i ON i.movie_id =
  m.movie_id
JOIN imdb_ratings r ON r.imdb_title_id =
  i.imdb_title_id
GROUP BY decade_released, m.year, g.genre
WITH ROLLUP
ORDER BY decade_released, m.year, g.genre;
```

NORMALIZED SCHEMA

```
SELECT CONCAT(FLOOR(m.year/10) * 10, 's') AS
  decade_released, m.year, g.genre,
  SUM(r.total_num_votes) AS total_votes,
  SUM(r.votes_10) AS votes_10, SUM(r.votes_9)
  AS votes_9, SUM(r.votes_8) AS votes_8,
  SUM(r.votes_7) AS votes_7, SUM(r.votes_6) AS
  votes_6, SUM(r.votes_5) AS votes_5,
  SUM(r.votes_4) AS votes_4, SUM(r.votes_3) AS
  votes_3, SUM(r.votes_2) AS votes_2,
  SUM(r.votes_1) AS votes_1
FROM fact_table f
JOIN movies m ON m.movie_id = f.movie_id
JOIN ratings r ON r.rating_id = f.rating_id
JOIN genres g ON g.genre_id = f.genre_id
GROUP BY decade_released, m.year, g.genre
WITH ROLLUP
ORDER BY decade_released, m.year, g.genre;
```

DENORMALIZED SCHEMA

OPTIMIZATION AND RESULTS



QUERY PROCESSING

Q1. For each genre, what is the total number of movie votes per year and decade?

Output

decade_released	year	genre	total_votes	votes_10	votes_9	votes_8	votes_7	votes_6	votes_5
NULL	NULL	NULL	60390538630	9063770405	9252537587	14024663182	12659453493	7400999751	3683119834
1890s	NULL	NULL	296433	31978	36868	80723	71746	36149	15923
1890s	1894	NULL	3056	184	138	474	1088	722	280
1890s	1894	Documentary	1528	92	69	237	544	361	140
1890s	1894	Short	1528	92	69	237	544	361	140
1890s	1895	NULL	16566	1128	822	2010	4572	4344	1950
1890s	1895	Drama	8283	564	411	1005	2286	2172	975
1890s	1895	Short	8283	564	411	1005	2286	2172	975
1890s	1896	NULL	25368	2280	1906	4457	6584	5074	2544
1890s	1896	Documentary	10334	784	523	1295	2250	2512	1466
1890s	1896	Drama	1036	42	39	74	222	316	200
1890s	1896	Short	13998	1454	1344	3088	4112	2246	878
1890s	1897	NULL	83962	9582	11298	21875	20238	10846	4782
1890s	1897	Comedy	9182	551	405	1347	2609	2115	953
1890s	1897	Crime	156	10	14	16	36	52	18
1890s	1897	Documentary	8382	571	601	1410	2109	1637	833
1890s	1897	Drama	32642	4196	5168	9481	7548	3380	1448
1890s	1897	Short	33600	4254	5110	9621	7936	3662	1530
1890s	1898	NULL	153283	17509	22476	51391	38376	13876	4820

Rows: 1992
Columns: 14



OPTIMIZATION AND RESULTS

Normalized

The diagram illustrates a normalized query plan for a join operation. At the top, a box labeled "GROUP" (representing a HashAggregate) receives input from "tmp_table, filesort" and sends output to "ORDER" (representing a Sort). The "ORDER" box also receives input from "query_block#1". Below this, a large box labeled "buffer_result" contains the main query plan. The plan starts with a "Full Table Scan" (red box) on "tmp_table" (399.75, 3.83K rows). This feeds into a "nested loop" join (diamond) with a "Unique Key Lookup" (green box) on "PRIMARY" (1247.1, 1 row). This join feeds into another "nested loop" join with a "Unique Key Lookup" on "PRIMARY" (5264.34, 1 row). This second join feeds into a third "nested loop" join with a "Non-Unique Key Lookup" (green box) on "PRIMARY" (2395.44, 2 rows). The final output of the plan is "10.28K rows".

Denormalized, with indexes

The diagram illustrates a denormalized database schema with indexes. It shows a 'tmp table' with three columns: 'Full Table Scan', 'idx_rating_id', and 'id'. The 'Full Table Scan' column has a value of 3715.15 and 35.14K rows. The 'idx_rating_id' column has a value of 2641403.17 and 214 rows, with a 'Non-Unique Key Lookup' label. The 'id' column has a value of 2641403.17 and 1 row, with a 'Unique Key Lookup' label. The 'tmp table' is connected to a 'buffer_result' table, which is connected to a 'GROUP' table, which is connected to an 'ORDER' table. The 'ORDER' table is connected to a 'Query Cost' table, which shows a cost of 76702.445 and a 'query_block #1'.

Denormalized, no indexes

Query cost: 1.880095358345e+11

query_block #1

ORDER

filesort

GROUP

tmp table, filesort

buffer_result

tmp table

2.35 21 rows

74337.71

hash join

741.34K rows

472093316530.21

hash join

47.21G rows

1880095358345.0

hash join

1880095358345 rows

248.02K rows

Full Table Scan

Full Table Scan

Full Table Scan

Full Table Scan

g

r

f

m

74335.36 35.30K rows

472135777.8 52M rows

18595443.2

Denormalized, composite indexes

The diagram illustrates a denormalized database structure with composite indexes. A query for 'ORDER' (query block #1) is executed. The execution plan shows a 'GATHER' step (tmp table result) feeding into an 'ORDER' step (result). The 'tmp table' is composed of three parts: a 'Full Table Scan' (3715.15, 25148 rows), a 'Non-Unique Key Lookup' (2641405.17, 214 rows), and a 'Unique Key Lookup' (2641405.17, 1 row). The 'tmp table result' is then processed by the 'ORDER' step, which has a 'result' step (792762.65, 72248 rows).



QUERY PROCESSING

Q2. For each year of each decade, what are the total number of male and female votes per genre?

```
SELECT CONCAT(FLOOR(m.year/10) * 10, 's') AS
  decade_released, m.year, g.genre,
  (SUM(r.votes_10) + SUM(r.votes_9) +
   SUM(r.votes_8) +
   SUM(r.votes_7) + SUM(r.votes_6) +
   SUM(r.votes_5) +
   SUM(r.votes_4) + SUM(r.votes_3) +
   SUM(r.votes_2) +
   SUM(r.votes_1)) AS total_num_votes,
  SUM(r.males_allages_votes) AS votes_male,
  SUM(r.females_allages_votes) AS votes_female
FROM movies m
JOIN genres g ON g.movie_id = m.movie_id
JOIN ijs2supplementary i ON i.movie_id =
  m.movie_id
JOIN imdb_ratings r ON r.imdb_title_id =
  i.imdb_title_id
GROUP BY decade_released, m.year, g.genre
WITH ROLLUP
ORDER BY decade_released, m.year, g.genre;
```

NORMALIZED SCHEMA

```
SELECT CONCAT(FLOOR(m.year/10) * 10, 's')
  AS decade_released, m.year, g.genre,
  SUM(r.total_num_votes) AS total_votes,
  SUM(r.males_allages_votes) AS votes_male,
  SUM(r.females_allages_votes) AS
    votes_female
FROM fact_table f
JOIN movies m ON m.movie_id = f.movie_id
JOIN genres g ON g.genre_id = f.genre_id
JOIN ratings r ON r.rating_id = f.rating_id
GROUP BY decade_released, m.year, g.genre
WITH ROLLUP
ORDER BY decade_released, m.year, g.genre;
```

DENORMALIZED SCHEMA



QUERY PROCESSING

Q2. For each year of each decade, what are the total number of male and female votes per genre?

Output

decade_released	year	genre	total_votes	votes_male	votes_female
NULL	NULL	NULL	60390538630	38694163614	8093541206
1890s	NULL	NULL	296433	183862	45811
1890s	1894	NULL	3056	1862	526
1890s	1894	Documentary	1528	931	263
1890s	1894	Short	1528	931	263
1890s	1895	NULL	16566	10836	2256
1890s	1895	Drama	8283	5418	1128
1890s	1895	Short	8283	5418	1128
1890s	1896	NULL	25368	14872	4524
1890s	1896	Documentary	10334	6131	1899

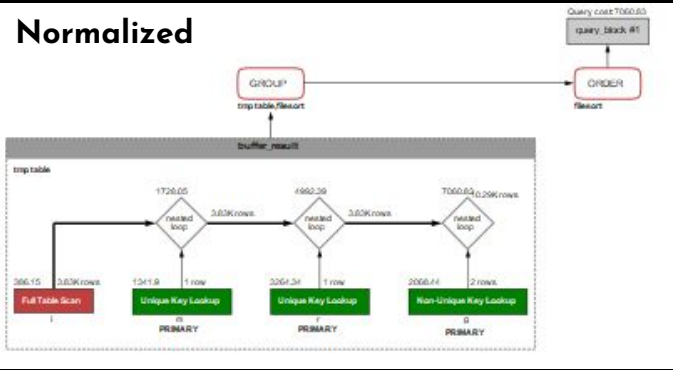
Rows: 1992

Columns: 6

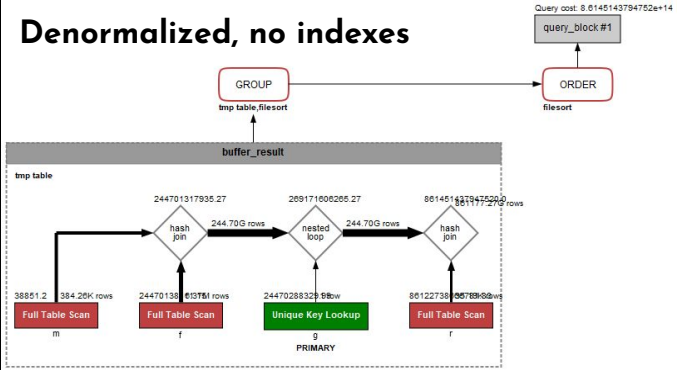


Q2. For each year of each decade, what are the total number of male and female votes per genre?

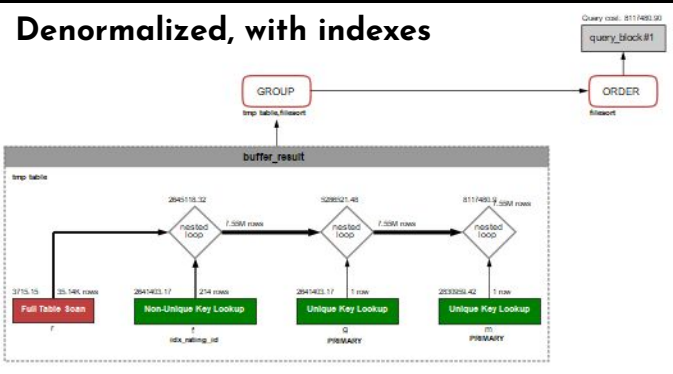
Query Execution Plans



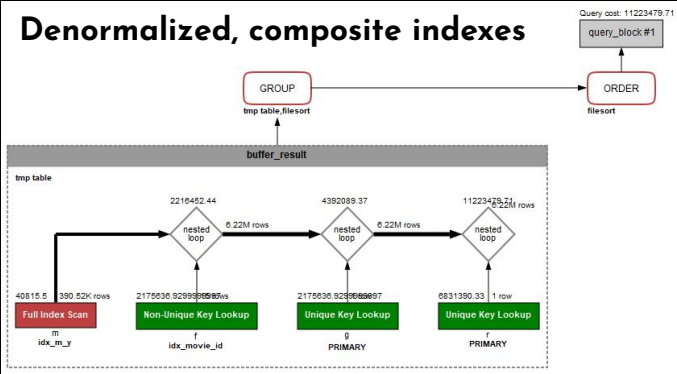
Execution Time: 12.5 seconds



Execution Time: 9.359 seconds



Execution Time: 12.656 seconds



Execution Time: 13.375 seconds

QUERY PROCESSING

Q3. Who are the top 15 actors with the most number of movie appearances for the action genre?

```
SELECT name, Num_roles
FROM (
  SELECT a.name, COUNT(m.movie_id)
    AS Num_roles,
    RANK() OVER
      (PARTITION BY g.genre
       ORDER BY COUNT(m.movie_id) DESC)
    AS Num_roles_rank
  FROM movies m
  JOIN genres g ON g.movie_id = m.movie_id
  JOIN roles i ON i.movie_id = m.movie_id
  JOIN actors a ON a.actor_id = i.actor_id
  WHERE g.genre = "Action"
  GROUP BY g.genre, a.name
) AS Role_rank
WHERE Num_roles_rank <= 15;
```

NORMALIZED SCHEMA

```
SELECT name, Num_roles
FROM (
  SELECT CONCAT(a.last_name, ", ",
    a.first_name) AS name, COUNT(m.movie_id)
    AS Num_roles,
    RANK() OVER
      (PARTITION BY g.genre
       ORDER BY COUNT(m.movie_id) DESC)
    AS Num_roles_rank
  FROM fact_table f
  JOIN genres g ON f.genre_id = g.genre_id
  JOIN actors a ON f.actor_id = a.actor_id
  JOIN movies m ON f.movie_id = m.movie_id
  WHERE g.genre = "Action"
  GROUP BY g.genre, a.last_name, a.first_name
) AS Role_rank
WHERE Num_roles_rank <= 15;
```

DENORMALIZED SCHEMA



QUERY PROCESSING

Q3. Who are the top 15 actors with the most number of movie appearances for the action genre?

Output

name	Num_roles
Almada, Mário	165
Chan, Jackie (I)	126
Reynoso, Jorge	126
Estrada, Joseph	115
Kapoor, Shakti (I)	105
Yuen, Biao	99
Puri, Amrish	96
García Jr., Eliazar	96
Park, No-shik (I)	94
Bernal, Agustín	90

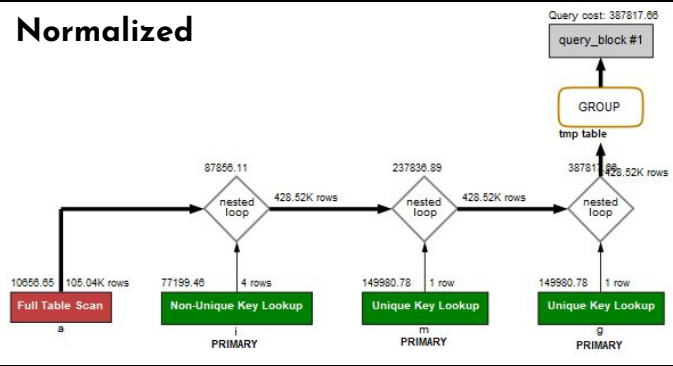
Rows: 15

Columns: 2

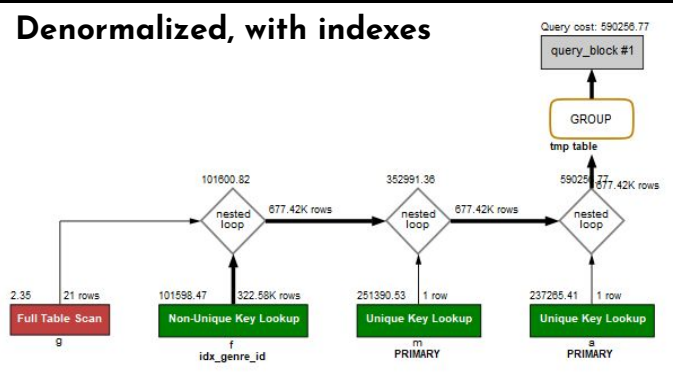


Q3. Who are the top 15 actors with the most number of movie appearances for the action genre?

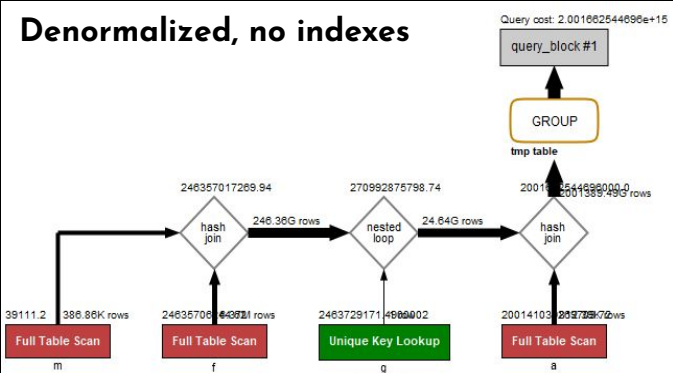
Query Execution Plans




Execution Time: 18.752 seconds



Execution Time: 2.547 seconds



Execution Time: 11.860 seconds



Execution Time: 1.391 seconds

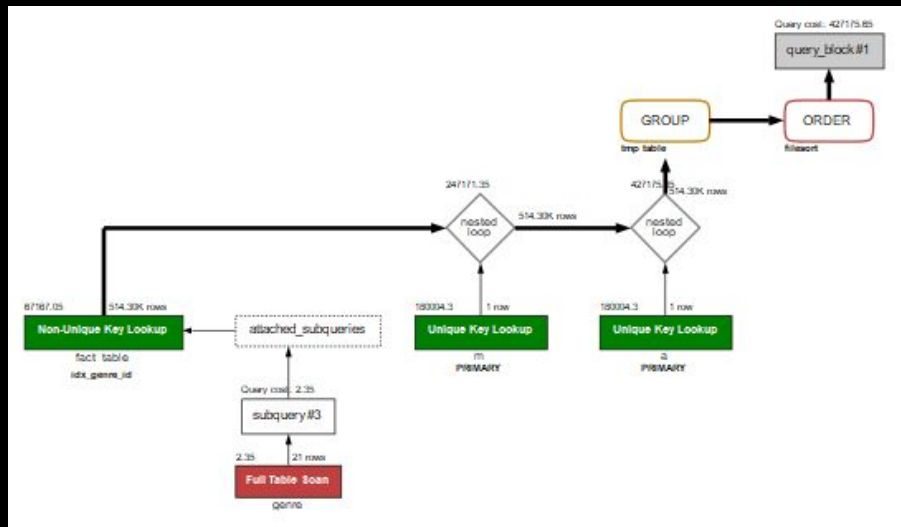
QUERY PROCESSING

Q3. Who are the top 15 actors with the most number of movie appearances for the action genre?

Optimized Query

```
SELECT CONCAT(a.last_name, ", ",  
             a.first_name)  
  AS name, COUNT(m.movie_id) AS  
  Num_roles  
FROM (  
  SELECT *  
  FROM fact_table  
  WHERE genre_id = (  
    SELECT genre_id  
    FROM genres  
    WHERE genre = "Action"))  
  AS f  
JOIN actors a ON f.actor_id = a.actor_id  
JOIN movies m ON f.movie_id = m.movie_id  
GROUP BY a.last_name, a.first_name  
ORDER BY Num_roles DESC  
LIMIT 15;
```

Query Execution Plan



Execution Time: 0.728 seconds



QUERY PROCESSING

Q4. For each country, what are the top 10 movie genres?

```
SELECT *
FROM (
  SELECT c.country, g.genre,
         COUNT(m.movie_id)
         AS Num_movies,
         RANK() OVER
         (PARTITION BY c.country
          ORDER BY COUNT(m.movie_id) DESC)
         AS Num_movies_rank
  FROM movies m
  JOIN countries c ON c.movie_id = m.movie_id
  JOIN genres g ON g.movie_id = m.movie_id
  WHERE c.country != ""
  GROUP BY c.country, g.genre
) AS Movies_rank
WHERE Num_movies_rank <= 10;
```

NORMALIZED SCHEMA

```
SELECT *
FROM (
  SELECT c.country, g.genre, COUNT(m.movie_id)
         AS Num_movies,
         RANK() OVER
         (PARTITION BY c.country
          ORDER BY COUNT(m.movie_id) DESC)
         AS Num_movies_rank
  FROM fact_table f
  JOIN countries c ON c.country_id = f.country_id
  JOIN genres g ON g.genre_id = f.genre_id
  JOIN movies m ON m.movie_id = f.movie_id
  WHERE c.country != ""
  GROUP BY c.country, g.genre
) AS Movies_rank
WHERE Num_movies_rank <= 10;
```

DENORMALIZED SCHEMA



QUERY PROCESSING

Q4. For each country, what are the top 10 movie genres?

Output

country	genre	Num_movies	Num_movies_rank
Afghanistan	Drama	14	1
Albania	Drama	39	1
Albania	Romance	31	2
Albania	Sci-Fi	31	2
Albania	Comedy	13	4
Albania	Family	6	5
Albania	Action	3	6
Albania	Documentary	1	7
Albania	Short	1	7
Algeria	Drama	187	1
Algeria	War	66	2
Algeria	Comedy	35	3

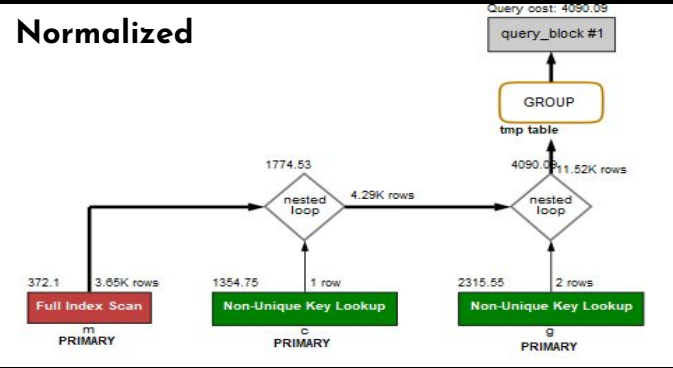
Rows: 1057

Columns: 4

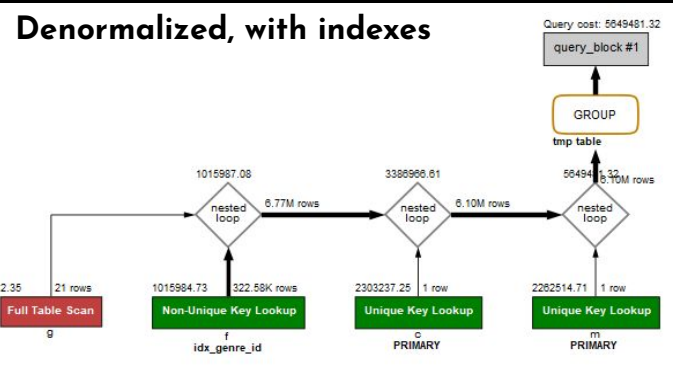


Q4. For each country, what are the top 10 movie genres?

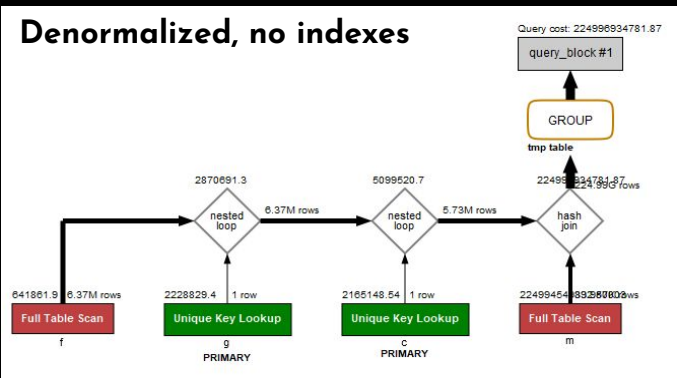
Query Execution Plans



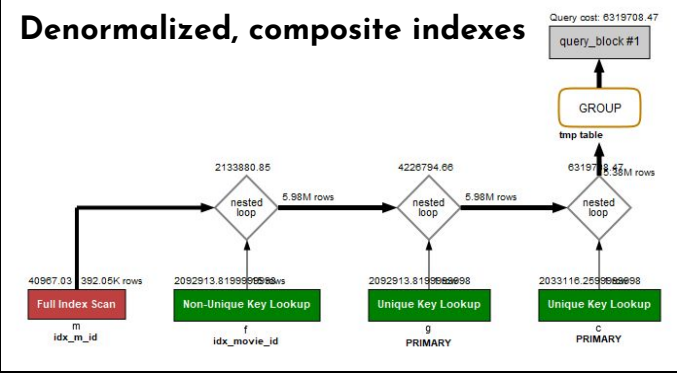
Execution Time: 25.682 seconds



Execution Time: 14.884 seconds



Execution Time: 21.129 seconds



Execution Time: 9.468 seconds

QUERY PROCESSING

Q5. How many horror movies are released in the USA for each year and decade?

```
SELECT CONCAT(FLOOR(year/10) * 10, 's')
  AS decade_released, m.year,
  COUNT(m.movie_id) AS
    num_movies
FROM fact_table f
JOIN genres g ON g.genre_id = f.genre_id
JOIN movies m ON m.movie_id = f.movie_id
JOIN countries c ON c.country_id =
  f.country_id
WHERE c.country = "USA"
  AND g.genre = "Horror"
GROUP BY decade_released, m.year
WITH ROLLUP;
```

NORMALIZED SCHEMA

```
SELECT CONCAT(FLOOR(year/10) * 10, 's')
  AS decade_released, m.year,
  COUNT(m.movie_id) AS
    num_movies
FROM fact_table f
JOIN genres g ON g.genre_id = f.genre_id
JOIN movies m ON m.movie_id = f.movie_id
JOIN countries c ON c.country_id =
  f.country_id
WHERE c.country = "USA"
  AND g.genre = "Horror"
GROUP BY decade_released, m.year
WITH ROLLUP;
```

DENORMALIZED SCHEMA

QUERY PROCESSING

Q5. How many horror movies are released in the USA for each year and decade?

Output

decade_released	year	num_movies
1990s	1998	1383
1990s	1999	1396
1990s	NULL	12076
2000s	2000	1604
2000s	2001	1300
2000s	2002	1929
2000s	2003	1762
2000s	2004	1787
2000s	2005	434
2000s	2006	5
2000s	NULL	8821
NULL	NULL	50221

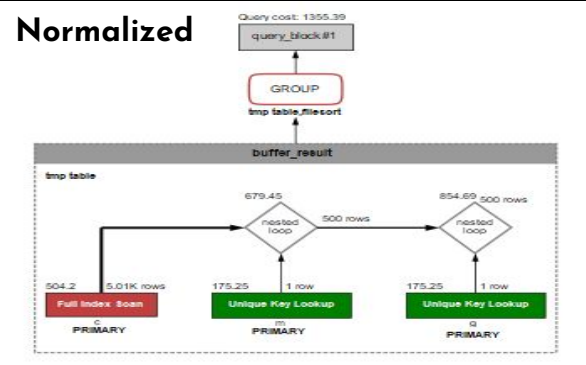
Rows: 99

Columns: 3

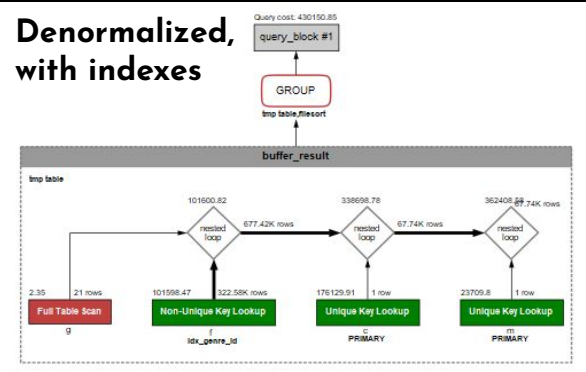


Q5. How many horror movies are released in the USA for each year and decade?

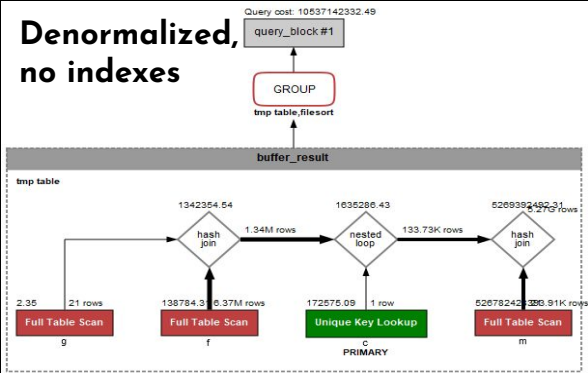
Query Execution Plans



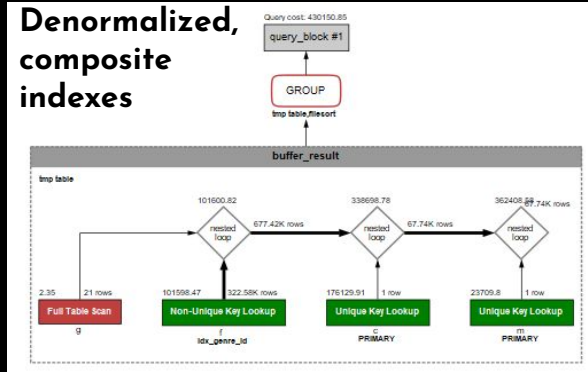
Execution Time: 31.154 seconds



Execution Time: 0.562 seconds



Execution Time: 2.860 seconds



Execution Time: 0.328 seconds

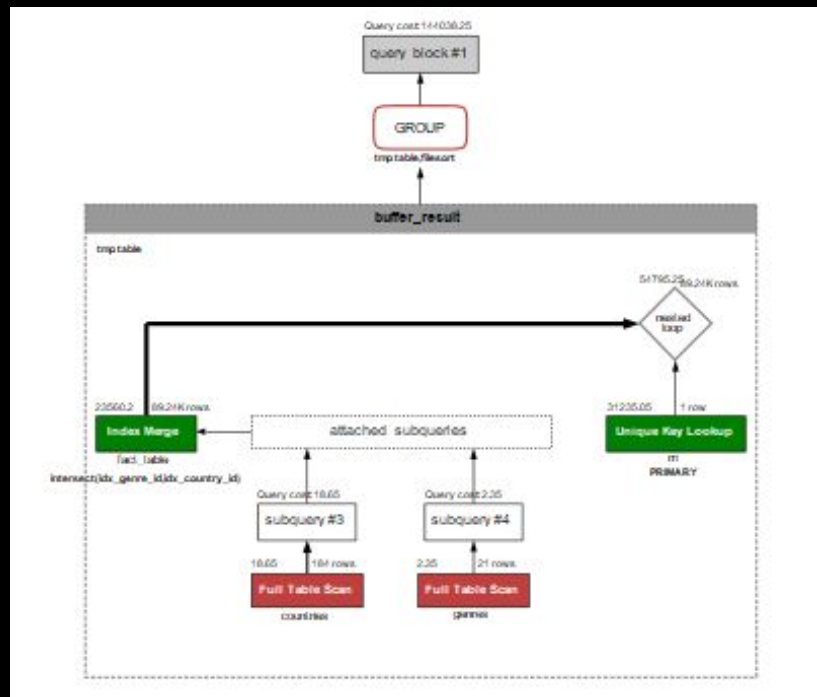
QUERY PROCESSING

Q5. How many horror movies are released in the USA for each year and decade?

Optimized Query

```
SELECT CONCAT(FLOOR(year/10) * 10, 's') AS  
    decade_released, m.year, COUNT(m.movie_id)  
    AS num_movies  
FROM (  
    SELECT *  
    FROM fact_table  
    WHERE country_id = (  
        SELECT country_id  
        FROM countries  
        WHERE country = "USA"  
    )  
    AND genre_id = (  
        SELECT genre_id  
        FROM genres  
        WHERE genre = "Horror"  
    )) AS f  
JOIN movies m ON m.movie_id = f.movie_id  
GROUP BY decade_released, m.year  
WITH ROLLUP;
```

Query Execution Plan



Execution Time: 0.152 seconds



QUERY PROCESSING

Q6. What are the production companies that have released over 10 movies with a rating above the average rating and a number of votes above the average total number of votes?

```
SELECT production_company, COUNT(imdb_title_id) AS num_movies
FROM (
  SELECT i.production_company, i.imdb_title_id
  FROM imdb_movies i
  WHERE i.imdb_title_id IN (
    SELECT r.imdb_title_id
    FROM imdb_ratings r
    WHERE r.mean_vote > (
      SELECT AVG(r.mean_vote)
      FROM imdb_ratings r
    )
  )
  AND imdb_title_id IN (
    SELECT i.imdb_title_id
    FROM imdb_movies i
    WHERE i.votes > (
      SELECT AVG(i.votes)
      FROM imdb_movies i
    )
  )
) AS selected_movies
GROUP BY production_company
HAVING num_movies > 10
ORDER BY num_movies DESC;
```

NORMALIZED SCHEMA

```
SELECT production_company, COUNT(movie_id) AS num_movies
FROM (
  SELECT m.movie_id, m.production_company,
    total_num_votes AS total
  FROM movies m
  JOIN fact_table ft ON ft.movie_id = m.movie_id
  JOIN ratings ra ON ra.rating_id = ft.rating_id
  WHERE m.production_company IS NOT NULL
  AND m.rank > (
    SELECT AVG(m.rank)
    FROM movies m
  )
  GROUP BY m.movie_id
  HAVING total > (
    SELECT AVG(total_num_votes)
    FROM (
      SELECT name, total_num_votes
      FROM fact_table f
      JOIN movies m ON m.movie_id = f.movie_id
      JOIN ratings r ON r.rating_id = f.rating_id
      JOIN genres g ON g.genre_id = f.genre_id
    ) as total_votes
  )
) AS selected_movies
GROUP BY production_company
HAVING num_movies > 10
ORDER BY num_movies DESC;
```

DENORMALIZED SCHEMA



QUERY PROCESSING

Q6. What are the production companies that have released over 10 movies with a rating above the average rating and a number of votes above the average total number of votes?

Output

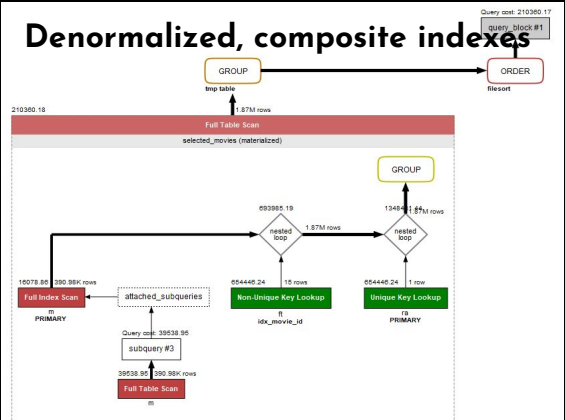
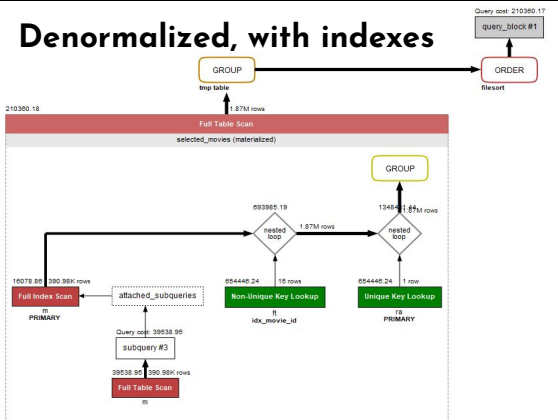
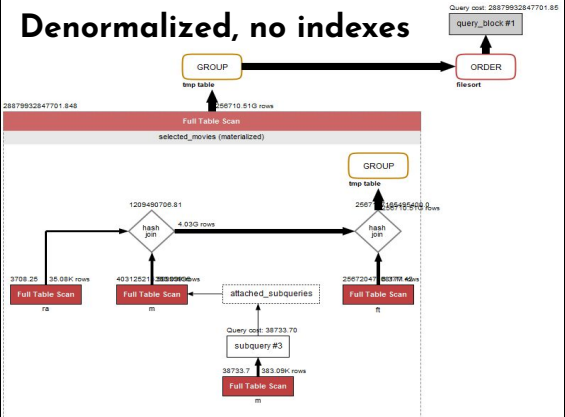
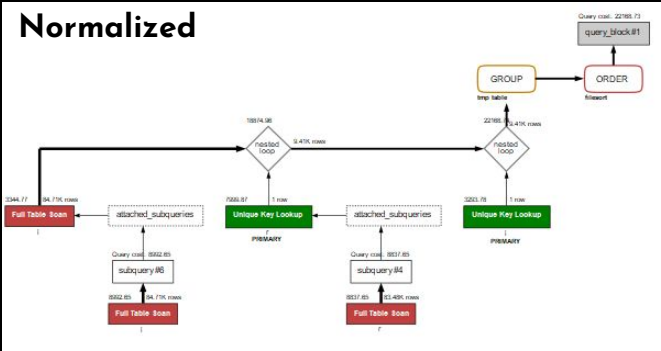
production_company ▼	num_movies
Warner Bros.	47
Universal Pictures	37
Twentieth Century Fox	31
Touchstone Pictures	18
Paramount Pictures	34
New Line Cinema	17
Metro-Goldwyn-Mayer (MGM)	13
Eon Productions	11
DreamWorks	11
Columbia Pictures	24

Rows: 10

Columns: 2



Q6. What are the production companies that have released over 10 movies with a rating above the average rating and a number of votes above the average total number of votes?



QUERY PROCESSING

Q7. For each of the top 6 directors with the highest gross earnings, what are their 15 highest-rated movies?

```
SELECT name, title, `rank`
FROM (
  SELECT m.title, topdir.name, topdir.gross,
         r.rank, ROW_NUMBER() OVER(PARTITION BY
         topdir.name ORDER BY `rank` DESC) AS
         director_movie_rank
  FROM (
    SELECT d.directorid, d.name, d.gross
    FROM directors d
    ORDER BY d.gross DESC
    LIMIT 6
  ) as topdir
  JOIN movies m
  JOIN ratings r ON r.movieid = m.movieid
  JOIN movies2directors md ON md.movieid =
    m.movieid AND md.directorid =
    topdir.directorid
  ORDER BY topdir.gross DESC, r.rank DESC
) AS topdirmovies
WHERE director_movie_rank <= 15;
```

NORMALIZED SCHEMA

```
SELECT full_name, gross, name, `rank`
FROM (
  SELECT full_name, m.name, gross, m.rank,
         ROW_NUMBER() OVER
         (PARTITION BY td.director_id ORDER BY
         `rank` DESC) AS Director_movie_rank
  FROM (
    SELECT d.director_id,
           CONCAT(CONCAT(d.last_name, ", "),
                  d.first_name) AS full_name, d.gross
    FROM directors d
    ORDER BY d.gross DESC
    LIMIT 6
  ) as td
  JOIN fact_table f ON td.director_id = f.director_id
  JOIN movies m ON m.movie_id = f.movie_id
  WHERE `rank` IS NOT NULL
  GROUP BY m.movie_id
  ORDER BY gross DESC, m.rank DESC
) AS top_director_movies
WHERE Director_movie_rank <= 15;
```

DENORMALIZED SCHEMA

OPTIMIZATION AND RESULTS



QUERY PROCESSING

Q7. For each of the top 6 directors with the highest gross earnings, what are their 15 highest-rated movies?

Output

	full_name	gross	name	rank
▶	Spielberg, Steven	5520276631	Escape to Nowhere	9.1
	Spielberg, Steven	5520276631	Firelight	9
	Spielberg, Steven	5520276631	Schindler's List	8.8
	Spielberg, Steven	5520276631	Raiders of the Lost Ark	8.7
	Spielberg, Steven	5520276631	Slipstream	8.6
	Spielberg, Steven	5520276631	The Last Gun	8.4
	Spielberg, Steven	5520276631	Amblin'	8.4
	Spielberg, Steven	5520276631	Saving Private Ryan	8.3
	Spielberg, Steven	5520276631	Jaws	8.2
	Spielberg, Steven	5520276631	Indiana Jones and the Last Crusade	8
	Spielberg, Steven	5520276631	Close Encounters of the Third Kind	7.8
	Spielberg, Steven	5520276631	E.T. the Extra-Terrestrial	7.8
	Spielberg, Steven	5520276631	Minority Report	7.8
	Spielberg, Steven	5520276631	Catch Me If You Can	7.7
	Spielberg, Steven	5520276631	The Color Purple	7.6
	Cameron, James...	3295280440	Aliens	8.2
	Cameron, James...	3295280440	Terminator 2: Judgment Day	8.1

Rows: 76

Columns: 4

OPTIMIZATION AND RESULTS



naïve nested-loop joins vs indexed nested-loop joins

B⁺-tree Lookup

$$O(\log n)$$

Time complexity

$$O(n)$$

Space complexity

Join algorithm costs

$$|r| \times b_s + b_r$$

naïve nested-loop join
block transfers

$$|r| + b_r$$

naïve nested-loop join
disk seeks

$$b_r \times (T_{transfer} + T_{seek}) + |r| \times c$$

indexed nested-loop join



- ✓ metadata of the local IMDb IJS database
was compared to the remote copy

```
SELECT TABLE_NAME, COLUMN_NAME,  
       ORDINAL_POSITION, IS_NULLABLE, COLUMN_TYPE,  
       COLLATION_NAME  
FROM INFORMATION_SCHEMA.COLUMNS  
WHERE table_schema = 'imdb_ijs';
```



- ✓ **CHECKSUM** used for IMDb IJS

```
CHECKSUM TABLE imdb_ijs.actors;
```

- ✓ **Boundary value analysis** for IMDb Supplementary

```
Maximum and minimum per column
```

- ✓ **Compare cardinality of table against CSV file**



Preprocessing Validation

✓ Movie Title Wrangling Validation (Comma)

```
SELECT COUNT(*) FROM imdb_star.movies  
WHERE `name` LIKE "%,";
```

✓ Movie Title Wrangling Validation (Article)

```
SELECT COUNT(*) FROM imdb_star.movies  
WHERE `name` LIKE "%, The";
```

✓ Genre Wrangling Validation

```
SELECT COUNT(*) FROM imdb_star.genres  
WHERE genre LIKE "%,%";
```

The results of these SQL statements should be 0 since the data have already been preprocessed



Preprocessing Validation

✓ Production Company Wrangling Validation (Brackets)

```
SELECT COUNT(*) FROM imdb_star.movies  
WHERE production_company LIKE "%[%]";
```

✓ Production Company Wrangling Validation (Parentheses)

```
SELECT COUNT(*) FROM imdb_star.movies  
WHERE production_company LIKE "%- (%)";
```

The results of these SQL statements should be 0 since the data have already been preprocessed



Performance Testing Results

Query	Normalized	Denormalized			
		No Indexes	With Indexes	With Composite Indexes	Revised Query
1	23.400	<u>9.344</u>	15.453	16.516	N/A
2	12.500	<u>9.359</u>	12.656	13.375	N/A
3	18.752	11.860	2.547	1.391	<u>0.728</u>
4	25.682	21.129	14.884	<u>9.468</u>	N/A
5	31.154	2.860	0.562	0.328	<u>0.152</u>
6	10.311	12.468	10.531	<u>8.762</u>	N/A
7	2.906	0.157	0.062	<u>0.032</u>	N/A



Deeper appreciation for
the forms of data and
issues arising from
synthesizing data from
different sources

Learning to integrate
external tools into the
pipeline



Optimization strategies
in query processing
perform differently on a
case-to-case basis





Moviegoers can use the results to examine the history of genres and predict film quality

Producers and directors can find prolific actors for a genre or examine the past production company projects



The results allow for insights and meaningful interpretations of the film industry



Optimizing the Dashboard

- ▶ Using Tableau's relations instead of joins
- ▶ Using Tableau's sets

TABLEAU
